

Lab Assignment 10: Convolutional Neural Networks with Python

Shiva Agarwal

May 16, 2024

Lab Assignment

Coding Question

You are required to implement a Convolutional Neural Network (CNN) using Python with Keras library for classifying images from the CIFAR-10 dataset. Follow the steps below:

1. Import the CIFAR-10 dataset and print training data-related information.
2. Import the necessary libraries for building and training the CNN model.
3. Specify the initial convolutional block of the CNN.
4. Add dropout layer for regularization.
5. Add another convolutional block for CNN.
6. Define dense layers that consume the feature array and produce a classification.
7. Compile and train the CNN model.
8. Evaluate the performance of the trained model.
9. Check for overfitting using appropriate visualization.
10. Make predictions using the trained model.
11. Print the confusion matrix to evaluate the model's performance.

Solution

Question 1: Importing Libraries

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from keras.callbacks import EarlyStopping
4 from keras.datasets import cifar10
5 from keras.models import Sequential
6 from keras.layers.core import Dense, Dropout, Flatten
7 from keras.layers import LeakyReLU
8 from keras.layers.convolutional import Conv2D
9 from keras.optimizers import Adam
10 from keras.layers.pooling import MaxPooling2D
11 from keras.utils import to_categorical
12 from sklearn.metrics import confusion_matrix
```

Question 2: Loading the CIFAR-10 Dataset

```
1 # Load CIFAR-10 dataset
2 (X_train, Y_train), (X_test, Y_test) = cifar10.load_data
   ()
3
4 # Print training data-related information
5 print("Shape of training data:", X_train.shape)
6 print("First 5 labels:", Y_train[:5])
```

Question 3: Building the CNN Model

```
1 # Build CNN model
2 classifier = Sequential()
3
4 classifier.add(Conv2D(32, kernel_size=(3, 3), padding='
   same', input_shape=(32, 32, 3)))
5 classifier.add(LeakyReLU(alpha=0.3))
6 classifier.add(Conv2D(64, padding='same', kernel_size=(3,
   3)))
7 classifier.add(LeakyReLU(alpha=0.3))
8 classifier.add(MaxPooling2D(pool_size=(2, 2)))
9 classifier.add(Dropout(0.25))
10
11 classifier.add(Conv2D(128, kernel_size=(3, 3)))
12 classifier.add(MaxPooling2D(pool_size=(2, 2)))
```

```

13 classifier.add(Conv2D(128, kernel_size=(3, 3)))
14 classifier.add(LeakyReLU(alpha=0.3))
15 classifier.add(MaxPooling2D(pool_size=(2, 2)))
16 classifier.add(Dropout(0.25))
17
18 classifier.add(Flatten())
19 classifier.add(Dense(1024))
20 classifier.add(LeakyReLU(alpha=0.3))
21 classifier.add(Dropout(0.5))
22 classifier.add(Dense(10, activation='softmax'))

```

Question 4: Compiling and Training the CNN Model

```

1 # Compile CNN model
2 classifier.compile(loss='categorical_crossentropy',
3                   optimizer=Adam(lr=0.0001, decay=1e-6),
4                   metrics=['accuracy'])
5
6 # Train CNN model
7 history = classifier.fit(X_train / 255.0, to_categorical(
8     Y_train),
9                             batch_size=128, shuffle=True,
10                             epochs=250,
11                             validation_data=(X_test / 255.0,
12                                             to_categorical(Y_test)),
13                             callbacks=[EarlyStopping(
14                                 min_delta=0.01, patience=4)])

```

Question 5: Evaluating the Trained Model

```

1 # Evaluate CNN model
2 scores = classifier.evaluate(X_test / 255.0,
3                             to_categorical(Y_test))
4 print('Loss:', scores[0])
5 print('Accuracy:', scores[1])
6
7 # Check for overfitting
8 plt.plot(history.history['loss'])
9 plt.plot(history.history['val_loss'])
10 plt.title('Model Loss')
11 plt.ylabel('Loss')
12 plt.xlabel('Epoch')
13 plt.legend(['Train', 'Test'], loc='upper left')

```

```
13 plt.show()
```

Question 6: Making Predictions

```
1 # Make predictions
2 sample_id = 108
3 plt.figure(figsize=(4, 4))
4 plt.imshow(X_test[sample_id])
5 plt.axis('off')
6 plt.show()
7
8 print("Predicted class:", np.argmax(classifier.predict(
    X_test[108].reshape((1,) + X_test[108].shape) / 255.))
    )
```

Question 7: Printing the Confusion Matrix

```
1 # Print confusion matrix
2 Y_preds = classifier.predict(X_test / 255.0)
3 cm = confusion_matrix(Y_test, Y_preds.argmax(axis=1))
4 print("Confusion Matrix:")
5 print(cm)
```